


# Mid-term review

# Q1



☐ 1. Multiple Choice: Bags, Queues, Stacks: You've been asked to program a bag in...


Points: 4

Question

You've been asked to program a bag in the knowledge that the number of elements in the bag will always be less than 10,000 and you have whatever memory you need. But the time to add an element must be constant. Also, the total time to iterate forwards or backward must be no worse than  $O(n)$ . With which data structure would you choose to implement the bag?

Answer

A. Linked list

 B. Array

C. Doubly-linked list

D. Hash table

# Q2



☐ 2. Multiple Choice: Bags, Queues, Stacks: You have been assigned to design a st... 

Points: 4

Question

You have been assigned to design a stack. It is essential that there is never a delay in pushing or popping the stack, that's to say these operations should be performed in  $O(1)$  time. Again, there are no memory constraints. Which data structure would be most suitable?

Answer


A. B-tree


B. Array

 C. Linked List

D. Doubly-linked list

# Q3



☐ 3. True/False: Cluster: You have developed an algorithm that ... 


Points: **3**

Question

You have developed an algorithm that takes  $O(n^2)$  time to solve a problem with  $n$  elements. The algorithm currently runs on a cluster of 10 nodes. Thus far,  $n$  has never exceeded 10,000. However, a new customer has requested a quote for running  $n = 20,000$ . You can spend \$20,000 on building the cluster out to 40 nodes or you can pay the same amount to a programmer who claims that he can refactor your algorithm to run in  $O(n \log n)$  time. You should start ordering the new hardware, right?

Answer

True

 False

# Q4

- My answer:
  - By building out the cluster to 40 nodes, we will be able to run the algorithm in the same time that we have been accustomed to previously.
  - But by employing the programmer to reduce the algorithm to  $O(n \lg n)$ , we can run the new algorithm on the existing cluster in less than 0.1% of the time that we have been taking up until now!!
- To score full marks:
  - Reasonable explanation *plus*
  - Some sort of *back-of-the-envelope* calculation

# Q5

<i><b>n</b></i>	<i><b>t</b></i>	<i><b>log n</b></i>	<i><b>n log n</b></i>	<i><b>n<sup>2</sup></b></i>	<i><b>log t</b></i>
8	12.4	3.00	24	64	3.63
16	32.2	4.00	64	256	5.01
32	79.4	5.00	160	1024	6.31
64	192.2	6.00	384	4096	7.59
128	443.7	7.00	896	16384	8.79

$\log t = 1.5 * \log n - 1$       i.e.  $t = \frac{1}{2} * n^{3/2}$       True relationship is  $t = n \log n$

# Q6

$n$	$t$	$\log n$	$n \log n$	$n^2$	$\log t$	$M$	$2^M$
8	12.4	3.00	24	64	3.63	3.5	11.31
16	32.2	4.00	64	256	5.01	5	32.00
32	79.4	5.00	160	1024	6.31	6.5	90.51
64	192.2	6.00	384	4096	7.59	8	256.00
128	443.7	7.00	896	16384	8.79	9.5	724.08
256	443.7	8.00	2048	65536	8.79	11	2,048.00

This uses the (incorrect) model but comes up with a reasonable answer.  
True model predicts 1024. The point is that with only four doublings, we can't be 100%

# Q7, Q8, Q9

- Some ambiguity in wording
- Q7: correct answer is  $n+1$
- Q8:  $\lg n$  or  $\lg (n+1)$
- Q9: question is incorrect so pretty much everyone gets full mark
- There are three regions into which new element can go:
  - Left (1): 1 comparison
  - Right (1): 2 comparisons
  - Middle ( $n-1$ ): 2 comparisons plus  $\log n$



# Q10, Q11


- **Prove** that sum of the numbers 1 through  $n$  is equal to  $n(n+1)/2$ 
  - By Induction:
    - Base case:  $n = 1$ ,  $S_1 = 1(2)/2 = 1$  (correct)
    - Inductive case: the increment  $\Delta$  from  $S_n$  to  $S_{n+1}$  should be  $n+1$
    - $\Delta$  equals  $\{(n+1)(n+2) - n(n+1)\}/2$
    - $\Delta = n+1$
    - QED -- 证毕 -- इति सिद्धम्
  - Or by arithmetic progression
    - $2 * S_n = (1 + 2 + \dots + n) + (n + n-1 + \dots + 1)$
    - $2 * S_n = n * (n+1)$
    - Therefore  $S_n = n * (n+1) / 2$
- Tilde notation:
  - $\sim n^2$  or  $\sim n^2/2$

# Q12

- Question:
  - Stirling's approximation for  $n!$ , the number of permutations of a list of  $n$  different objects, expressed in entropy, as bits, is:  $n \lg n / e + \ln(2 \pi n) / 2$ . Express this in  $\sim$  ("tilde") notation.
- Answer:
  - $\sim n \lg n$
- Notes: it's important not to include any other terms in the tilde notation.

# Q13

1

☐ 13. Multiple Choice: Simple Sorts: Selection sort and Insertion sort are... 

Points: 5

Question

Selection sort and Insertion sort are both  $O(n^2)$  algorithms so they are only suitable for relatively small  $n$ . Which of the following are true?  
  
a. Insertion sort does half as many comparisons, generally speaking, as Selection sort;  
b. Insertion sort can reduce the time for exchanging elements by moving blocks of elements;  
c. Insertion sort is linear when the list is already sorted--unfortunately, this is not true for Selection sort.

Answer

1. a only


2. b only

3. c only

4. a and b

5. b and c

6. a and c

 7. all of the above

# Q14

- You are required to implement a method for the storage of up to 1 million elements. Each element has a key which represents a total ordering. It is desired to be able to select any element according to its key. However, you expect that the total number of different possible keys is somewhat greater than 4 billion (the number of possible *int* values). You don't want to search for these elements by traversing the list and comparing keys. But obviously you don't want to assign array storage to have one open slot for every possible key value! The way to do this is to implement a hash table. When you add an element to the table, you compute its *hashCode* and map that (typically by shifting bits to the right) into an index value which points to an element of the table. When it's possible (as it usually is) for many values to map to the same element (called a *collision*) there are several schemes to deal with that. The simplest is to use the next higher empty slot.
- However, for lookup (in the table) and comparison purposes, you will use the *hashCode* as a surrogate for the key itself. The *hashCode* is monotonically increasing with the key. That's to say that if  $hashCode(x) > hashCode(y)$ , then  $x > y$ . However, since the *hashCode* is essentially a 32-bit digest of the key, the *hashCode* itself does not qualify as a total order for the elements. That's because it is possible that  $x > y$  or  $x < y$  while, at the same time,  $hashCode(x) = hashCode(y)$ .
- If there are 100 million possible different values in the domain of the values (that's to say there are 100 million possible values of the key), what is the (approximate) probability that two different elements will have the same hash key (assume that the *hashCode* is uniformly distributed over the domain)?
- Answer: 1 in 43 ( $2^{32} / 100M$ )

# Q15

```
// Sort the arrays indices and hashes by comparing hashes
private void insertionSort(int n, int[] indices, int[] hashes) {
    // TODO implement me (8 points)
    for( int i=1; i< n; i++ ) {
        for( int j=i; j>0 && (hashes[j]<hashes[j-1]); j-- ) {
            exchange(hashes, j);
            exchange(indices, j);
        }
    }
}

// Verify that the arrays are in true order according to natural ordering on X
private void verify(int n, int[] indices, int[] hashes, List<X> a) {
    // TODO implement me (8 points)
    for(int i=0; i<n-1; i++) {
        if( hashes[i]==hashes[i+1] ) {
            if( a.get(indices[i]).compareTo( a.get(indices[i+1])) <0 )
                continue;
            else
                exchange(indices,i);
        }
    }
}
```